

Reinforcement Learning

Hyoung Joon, Kim

July 15, 2026

Contents

1	Introduction	4
I	Algorithms	5
2	Dynamic Programming	6
2.1	Introduction	6
2.2	Value Iteration	6
2.2.1	Main Idea	6
2.2.2	Pseudocode	7
2.2.3	Convergence	7
2.3	Policy Iteration	7
2.3.1	Main Idea	7
2.3.2	Pseudocode	8
2.3.3	Convergence	8
2.4	Truncated Policy Iteration	8
2.4.1	Main Idea	8
2.4.2	Pseudocode	9
2.4.3	Convergence	9
3	Monte Carlo Methods	10
3.1	Introduction	10
3.2	Monte Carlo Basic	10
3.2.1	Main Idea	10
3.2.2	Pseudocode	11
3.2.3	Convergence	11
3.2.4	Pros and Cons	11
3.3	Monte Carlo with Exploring Starts	11
3.3.1	Main Idea	11
3.3.2	Pseudocode	12
3.3.3	Implementation Details	12
3.3.4	Pros and Cons	12
4	Temporal Difference Methods	14
5	Policy-based Methods	15
5.1	Introduction	15
5.2	REINFORCE	15

5.2.1	Main Idea	15
5.2.2	Objective Function	15
5.2.3	Pseudocode	16
5.2.4	Implementation Notes	16
5.2.5	Pros and Cons	16
6	Model-based Methods	17
7	Comparisons of Algorithms	18
II	Theories	19
8	Basic Concepts	20
8.1	Introduction	20
8.2	State, Action, Policy and Reward	20
8.2.1	State	20
8.2.2	Action	20
8.2.3	Policy	21
8.2.4	Reward	21
8.3	Trajectories, Returns and Episodes	21
8.3.1	Trajectory	21
8.3.2	Returns	21
8.3.3	Episodes	22
8.4	Markov Decision Process	22
8.4.1	Definition	22
9	Bellman Equation	23
9.1	Introduction	23
9.2	Bellman Equation	23
9.2.1	State-value Function	23
9.2.2	Bellman Equation for State-Value Function	23
9.2.3	Matrix-Vector form of Bellman Equation	24
9.3	Solving The Bellman Equation	25
9.3.1	Closed-form Solution	25
9.3.2	Iterative Solution	26
9.4	Bellman Equation for Action-Value Function	26
9.4.1	Action Value	26
9.4.2	Bellman Equation for Action-Value Function	27
10	Bellman Optimality Equation	28
10.1	Introduction	28
10.2	Optimal Policy	28
10.3	Bellman Optimality Equation	29
10.4	Solving The Bellman Optimality Equation	29
10.4.1	Contraction Mapping Theorem	29
10.4.2	Contraction Property of Bellman Optimality Equation	31

10.5 Optimal Policy, Optimal State Value and Bellman Optimality Equation	32
10.6 Factors that Influence Optimal Policies	33
10.6.1 Reward	33
10.6.2 Discount Rate	33
III Appendices	34
11 Probability Theory	35
11.1 Introduction	35
11.2 Law of Total Expectation	35
12 Linear Algebra	36
12.1 Introduction	36
12.2 Geshgorin Circle Theorem	36

Chapter 1

Introduction

This document is a set of notes of algorithms of reinforcement learning and theories of reinforcement learning. Following books were mainly used for studying algorithms and theories.

- [\[Sut20\]](#);
- [\[Gra20\]](#);
- [\[Zha25\]](#);

Part I

Algorithms

Chapter 2

Dynamic Programming

2.1 Introduction

This chapter discuss and dynamic programming algorithms for finding optimal policies. Since these are dynamic programming algorithms, a perfect system model is required; this mean that it is rarely used in practice, but these algorithms important foundations of the model-free algorithms. Three algorithms will be presented;

- Value Iteration: The algorithms suggested by the contraction mapping theorem [10.5.0.1](#);
- Policy Iteration:
- Truncated Policy Iteration: The unified algorithm which includes value iteration and policy iteration as its special cases.

2.2 Value Iteration

2.2.1 Main Idea

The main idea of **value iteration** algorithms is the algorithms suggested in Theorem [10.5.0.1](#). In particular, the algorithms is

$$v_{k+1} = \max_{\pi \in \Pi} (r_{\pi} + \gamma P_{\pi} v_k), \quad k = 0, 1, 2, \dots \quad (2.1)$$

Breaking down, this iterative algorithm has two steps;

1. *Policy Update*: The algorithm finds a policy π_{k+1} that solves following optimization problem;

$$\arg \max_{\pi} (r_{\pi} + \gamma P_{\pi} v_k) \quad (2.2)$$

2. *Value Update*: With the policy π_{k+1} , the algorithm calculates a new state value function v_{k+1} ;

$$v_{k+1} = r_{\pi_{k+1}} + \gamma P_{\pi_{k+1}} v_k \quad (2.3)$$

and this v_{k+1} is used in next iteration.

2.2.2 Pseudocode

Algorithm 1: Value Iteration Algorithm

Input : $p(s', r|s, a)$ the dynamics of the system
 v initial guess

Output: v^*, π^* optimal state value and optimal policy

```

1  $\pi \leftarrow$  greedy policy derived from  $v$ 
2 repeat
3   for  $s \in \mathcal{S}$  do
4     for  $a \in \mathcal{A}(s)$  do
5        $q(s, a) \leftarrow \sum_r p(r|s, a)r + \gamma \sum_{s'} p(s'|s, a)v(s')$ 
        // policy update
6        $\pi(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_a q(s, a) \\ 0 & \text{otherwise} \end{cases}$ 
        // value update
7        $v(s) = \max_a q(s, a)$ 
8 until convergence of  $v$ 
9 return  $v, \pi$ 
```

Note that the value update of Algorithm 1 is simply substituting the original state values with maximum action values since we are using greedy policy.

2.2.3 Convergence

The convergence can be proved with Theorem 10.4.1.1 and Theorem 10.4.2.1.

2.3 Policy Iteration

2.3.1 Main Idea

The **policy iteration** is based on the policy improvement theorem. Like value iteration, the policy iteration consists of two steps:

1. *Policy Evaluation*: For a given policy π_k , we calculate the corresponding state values by solving following Bellman equation:

$$v_{\pi_k} = r_{\pi_k} + \gamma P_{\pi_k} v_{\pi_k} \quad (2.4)$$

2. *Policy Improvement*: Using the computed state value v_k , we obtain the improved policy as

$$\pi_{k+1} = \arg \max_{\pi} (r_{\pi} + \gamma P_{\pi} v_{\pi_k}) \quad (2.5)$$

and here is where the policy improvement theorem is applied.

The difference is that the value iteration is that instead of taking a state value and computing the policy from it, the policy iteration takes a policy and compute the value corresponding to the policy.

Also, note that policy iteration has another iterative algorithm inside it: the policy evaluation uses iterative algorithm for solving the Bellman equation $v_{\pi_k} = r_{\pi_k} + \gamma P_{\pi_k} v_{\pi_k}$.

2.3.2 Pseudocode

Algorithm 2: Policy Iteration Algorithm

Input : $p(s', r|s, a)$ the dynamics of the system
 π initial guess

Output: v^*, π^* optimal state value and optimal policy

```

1 repeat
    // policy evaluation
2    $v' \leftarrow$  arbitrary initial guess
3   repeat
4     for each state  $s \in \mathcal{S}$  do
5        $v'(s) \leftarrow \sum_a \pi(a|s) \left[ \sum_r p(r|s, a)r + \gamma \sum_{s'} p(s'|s, a)v'(s') \right]$ 
6   until convergence of  $v'$ 
7    $v \leftarrow v'$ 
    // policy improvement
8   for each state  $s \in \mathcal{S}$  do
9     for each action  $a \in \mathcal{A}(s)$  do
10       $q(s, a) \leftarrow \sum_r p(r|s, a)r + \gamma \sum_{s'} p(s'|s, a)v(s')$ 
11       $a^* \leftarrow \arg \max_a q(s, a)$ 
12       $\pi(a|s) \leftarrow \begin{cases} 1 & \text{if } a = a^* \\ 0 & \text{otherwise} \end{cases}$ 
13 until convergence of  $v$ 
14 return  $v, \pi$ 
```

2.3.3 Convergence

TODO: convergence proof of policy iteration theorem. This will be done after studying the operator version of Bellman Equation.

2.4 Truncated Policy Iteration

2.4.1 Main Idea

The **truncated policy iteration** is a more general algorithm including value iteration and policy iteration as its special cases. The algorithm has only one difference with policy iteration: in the policy evaluation step, the truncated policy iteration iterates only finite times, even when the convergence condition is not met. Truncated policy iteration algorithm is one of the instances of **general policy iteration**, which is an abstract concept of interleaving evaluation and improvement.

The perspective of looking at the truncated policy iteration as the generalization of policy iteration and value iteration can

be derived from the similiarity between value iteration and policy iteration. With the assumption that $v_{\pi}^{(0)} = v_0 = v_{\pi_0}$,

$$v_{\pi_1}^{(0)} = v_0 \quad (2.6)$$

$$\text{value iteration} \leftarrow v_1 \leftarrow v_{\pi_1}^{(1)} = r_{\pi_1} + \gamma P_{\pi_1} v_{\pi_1}^{(0)} \quad (2.7)$$

$$v_{\pi_1}^{(2)} = r_{\pi_1} + \gamma P_{\pi_1} v_{\pi_1}^{(1)} \quad (2.8)$$

$$v_{\pi_1}^{(3)} = r_{\pi_1} + \gamma P_{\pi_1} v_{\pi_1}^{(2)} \quad (2.9)$$

$$\vdots \quad (2.10)$$

$$\text{truncated policy iteration} \leftarrow \bar{v}_1 \leftarrow v_{\pi_1}^{(j+1)} = r_{\pi_1} + \gamma P_{\pi_1} v_{\pi_1}^{(j)} \quad (2.11)$$

$$\vdots \quad (2.12)$$

$$\text{policy iteration} \leftarrow v_{\pi_1} \leftarrow v_{\pi_1}^{(\infty)} = r_{\pi_1} + \gamma P_{\pi_1} v_{\pi_1}^{(\infty)} \quad (2.13)$$

$$(2.14)$$

While value iteration computes only once and update the state value and policy iteration iterates until convergence and use it as state value of policy, the truncated policy iteration iterates only finite number of times.

2.4.2 Pseudocode

Algorithm 3: Truncated Policy Iteration Algorithm

Input : $p(s', r|s, a)$ the dynamics of the system

π initial guess

Output: v^*, π^* optimal state value and optimal policy

```

1 repeat
    // policy evaluation
2    $v' \leftarrow$  arbitrary initial guess
3   for  $k = 0, 1, \dots, k_{\text{truncate}} - 1$  do
4       for each state  $s \in \mathcal{S}$  do
5            $v^{(k+1)}(s) \leftarrow \sum_a \pi(a|s) \left[ \sum_r p(r|s, a)r + \gamma \sum_{s'} p(s'|s, a)v^{(k)}(s') \right]$ 
6    $v \leftarrow v^{(k_{\text{truncate}})}$ 
    // policy improvement
7   for each state  $s \in \mathcal{S}$  do
8       for each action  $a \in \mathcal{A}(s)$  do
9            $q(s, a) \leftarrow \sum_r p(r|s, a)r + \gamma \sum_{s'} p(s'|s, a)v(s')$ 
10       $a^* \leftarrow \arg \max_a q(s, a)$ 
11       $\pi(a|s) \leftarrow \begin{cases} 1 & \text{if } a = a^* \\ 0 & \text{otherwise} \end{cases}$ 
12 until convergence of  $v$ 
13 return  $v, \pi$ 

```

2.4.3 Convergence

TODO: Convergence proof of truncated policy iteration.

Chapter 3

Monte Carlo Methods

3.1 Introduction

This chapter introduces a **Monte Carlo methods** for reinforcement learning. These algorithms all starts from trying to obtain the optimal policy by solving the Bellman optimality equation. The difference of pure dynamic programming algorithms and Monte Carlo methods is that the latter do not require a perfect dynamics of the model, i.e., it is *model-free*.

Then, how do we get the optimal policy without a perfect model? We use the data that the model gives us. Using the data, we can convert the policy iteration algorithm 2 to the model-free version. Specifically, we estimate the action value $q(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$ using Monte Carlo sampling.

3.2 Monte Carlo Basic

3.2.1 Main Idea

The **Monte Carlo basic** is the simplest Monte Carlo method algorithm. Suppose that we run the agent starting from state-action pair (s, a) , following policy π , and let $g_\pi^{(i)}(s, a)$ be the return of the episode. Then, we can approximate the action value $q_\pi(s, a)$ as following:

$$q_\pi(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a] \approx \frac{1}{n} \sum_{i=1}^n g_\pi^{(i)}(s, a). \quad (3.1)$$

If the number n is sufficiently large, the approximation will be sufficiently accurate. This can be guaranteed due to the law of large numbers.

3.2.2 Pseudocode

Algorithm 4: Monte Carlo Basic

Input : π_0 initial guess
Output: π^* optimal policy

```
1 repeat
2   foreach  $s \in \mathcal{S}$  do
3     foreach  $a \in \mathcal{A}(s)$  do
4       // policy evaluation
5       collect sufficiently many episodes starting from  $(s, a)$ , following policy  $\pi_k$ 
6        $q_{\pi_k}(s, a) \leftarrow$  the average return of all episodes starting from  $(s, a)$ 
7     // policy improvement
8      $\pi_{k+1}(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_{a'} q_{\pi_k}(s, a') \\ 0 & \text{otherwise} \end{cases}$ 
9   until convergence of  $\pi$ 
10  return  $\pi$ 
```

3.2.3 Convergence

3.2.4 Pros and Cons

Cons:

- Very inefficient sample usage. Due to this property, MC-Basic is not used in practical. To deal with this, we use MC with exploring start, which is presented in next section.

3.3 Monte Carlo with Exploring Starts

3.3.1 Main Idea

The MC-Basic algorithms is not sample-efficient, since it only updates only q value of one state-action pair each episode. To deal with this low efficiency, we use all the *visit* to the state-action pair in each episode.

Suppose that the agent has gone through following trajectory:

$$s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} s_4 \xrightarrow{a_4} \dots \quad (3.2)$$

To increase the sample efficiency, we use the visits to the state-action pairs in the trajectory. That is, using the above episode, we not only use (s_1, a_1) for estimating $q(s_1, a_1)$ but use (s_2, a_2) for estimating $q(s_2, a_2)$, and so on. There are two different ways to do this:

- The first way is to only use the state-action pairs when it is first encountered. This is called *first-visit method*;
- The second way is to use all the visit to the state-action pairs. This is called *every-visit method*.

Also, we update the estimate of q values after each episode ends. This falls into the scope of *generalized policy iteration*, since we are using the value approximation even it is not so accurate.

3.3.2 Pseudocode

Algorithm 5: Monte Carlo with Exploring Starts (every-visit)

Input : π_0 initial policy
Output: π^* optimal policy

```

// initialize
1 foreach valid state-action pair  $(s, a)$  do
2    $\pi \leftarrow \pi_0$ 
3    $q(s, a) \leftarrow$  initial value
   // holds the sum of returns of (sub)episode starting from  $(s, a)$ 
4    $\text{Ret}(s, a) \leftarrow 0$ 
   // counts the number of visits of  $(s, a)$ 
5    $N(s, a) \leftarrow 0$ 
6 repeat
7   select a starting valid state-action pair  $(s, a)$ , ensuring that all the pairs can be possibly selected
8   generate an episode starting from  $(s, a)$ , following current policy  $\pi$ :  $s_0, a_0, r_1, s_1, a_1, \dots, s_{T-1}, a_{T-1}, r_T$ 
9    $G \leftarrow 0$ 
10  for  $t = T - 1, T - 2, \dots, 0$  do
11     $G \leftarrow r_{t+1} + \gamma G$ 
12     $\text{Ret}(s, a) \leftarrow \text{Ret}(s, a) + G$ 
13     $N(s, a) \leftarrow N(s, a) + 1$ 
    // policy evaluation
14     $q(s_t, a_t) \leftarrow \frac{\text{Ret}(s, a)}{N(s, a)}$ 
    // policy improvement
15     $\pi_{k+1}(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_{a'} q_{\pi_k}(s, a') \\ 0 & \text{otherwise} \end{cases}$ 
16 until convergence

```

3.3.3 Implementation Details

The exploring-start condition can be achieved by random start state.

3.3.4 Pros and Cons

Pros:

- Better sample efficiency than MC-Basic;
- All state-action pairs can be visited due to the exploring-start, and this makes the estimation more accurate.

Cons:

- Only applicable to episodic tasks, which is a common drawback for all MC methods. This property leads to following problem, too;
- High variance since the return holds all the noises and randomnesses, and the return is used for update;
- If the exploring start condition cannot be satisfied, the MC with Exploring Starts cannot be used. This is common in real-world problems. To deal with this, another algorithms without exploring starts will be introduced next section;

Chapter 4

Temporal Difference Methods

Chapter 5

Policy-based Methods

5.1 Introduction

Policy gradient method, unlike value-based methods, tries to predict the optimal policy itself, instead of solving the Bellman equation. It has many advantages over value-based methods; It is more efficient in handling large state-action spaces, and has stronger generalization abilities.

5.2 REINFORCE

5.2.1 Main Idea

REINFORCE is the most basic form of policy gradient method. It's the vanilla policy gradient method.

5.2.2 Objective Function

REINFORCE uses following objective function;

$$J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}}[G_t] = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \gamma^t r_{t+1} \right]. \quad (5.1)$$

We use some tricks to compute the gradient of $J(\pi_{\theta})$ Then we get

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} G_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]. \quad (5.2)$$

Now we can use this to update the parameter θ . Since we cannot compute the exact expectation, we use Monte Carlo sampling to approximate $\nabla_{\theta} J(\pi_{\theta})$ as following:

$$\nabla_{\theta} J(\pi_{\theta}) \approx \sum_{t=0}^{T-1} G_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \quad (5.3)$$

5.2.3 Pseudocode

Algorithm 6: REINFORCE

Input : θ initial parameter
 γ discount rate
 α learning rate
Output: π_θ learned policy

```
1 initialize policy  $\pi_\theta$ 
2 repeat
3   generate episode  $\{s_0, a_0, r_1, \dots, s_{T-1}, a_{T-1}, r_T\}$  following  $\pi_\theta$ 
4   for  $t = 0, 1, \dots, T - 1$  do
5     // compute return-to-go
6      $G_t = \sum_{k=t+1}^T \gamma^{k-t-1} r_k$ 
7     // policy update
8      $\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(a_t | s_t) G_t$ 
9 until convergence
10 return  $\pi_\theta$ 
```

5.2.4 Implementation Notes

- When computing return-to-go, we can iterate backward from the terminal time to start, since $G_t \leftarrow \gamma G_{t+1} + r_{t+1}$ where $t = 0, 1, 2, \dots, T - 1$.

5.2.5 Pros and Cons

Pros:

- No need for model;
- Can cover both discrete and continuous actions;

Cons:

- Large variance;
- May converge to suboptimal policy;

Chapter 6

Model-based Methods

Chapter 7

Comparisons of Algorithms

Part II

Theories

Chapter 8

Basic Concepts

8.1 Introduction

In this section, we discuss common concepts that are used frequently in reinforcement learning.

8.2 State, Action, Policy and Reward

8.2.1 State

The **state** is the status of the agent with respect to the environment. The set of all possible states, the **state space**, is denoted as $\mathcal{S}$.

8.2.2 Action

The **action** is what agent can do, when a state is given. By choosing the action and taking it, the agent moves to the next state, which is called **state transition**. The set of all possible actions in state s , **action space** for state s , is denoted as $\mathcal{A}(s)$.

8.2.2.1 Transition Probability

Note that the transition can be stochastic. That is, the same choice of action in the same state may not lead to equivalent next state. We denote this probability with conditional probability as following:

$$p(s'|s, a) \tag{8.1}$$

where $s', s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$.

8.2.3 Policy

The **policy** tells the agent which actions to take at every state. The policy can be deterministic or stochastic. We denote the probability of taking action a under state s following policy π as $\pi(a|s)$.

8.2.4 Reward

Reward is one of the most unique concept of reinforcement learning. After the execution of action, the agent receives a reward, denoted as r . The agent's goal is to maximize the total reward that the agent receives. Due to this goal, the reward can encourage or discourage the agent's choice of selecting an action under certain state. Again, the reward can be stochastic, too.

One question that follows naturally is: Can't we just pick the actions that gives the most reward at every state? If this was correct, the reinforcement learning may not have existed. Since the reward is just an immediate information of the act, choosing only the action with highest reward is not the correct answer, most of the time. There can be a better choice in the long run.

8.3 Trajectories, Returns and Episodes

8.3.1 Trajectory

The **trajectory** is the footsteps of the agent. More formally, suppose that the agent is on the initial state s_0 and took action a_0 . The agent receives reward r_1 from environment and moves to state s_1 . Again, the agent take an action a_1 , receive reward r_1 and move to state s_2 . Repeating this, we get the trajectory of an agent:

$$s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T. \quad (8.2)$$

8.3.2 Returns

Given the trajectory $\tau = \{s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\}$, we can calculate the **return** G_t , of the trajectory. The naive definition of the return, called **total rewards** is the sum

$$G_t = \sum_{k=t+1}^T r_k. \quad (8.3)$$

The point is that we do not know if the terminal time T is finite or not. If T is not finite, than the cummulative reward [8.3](#) diverges. For this, we introduce the **discounted return**, defined as

$$G_t = \sum_{k=t+1}^{\infty} \gamma^{k-t-1} r_k \quad (8.4)$$

where $\gamma \in [0, 1]$ is called **discount rate**. This works for for finite trajectory too, if we let $r_k = 0$ for $k > T$.

8.3.3 Episodes

The **episode** is the finite trajectory itself, like 8.2. The episode may stop at the states called **terminal states**. The tasks with episodes are called **episodic tasks**. Unlike episodic tasks, there are **continuing tasks**, which does not ends. These two can be expressed in unified mathematical manner. For this, in the episodic task, we must consider the agent's behavior after the terminal state: to stay in that state forever, or to take action like normal states. In this note, we follow the latter, like [Zha25].

8.4 Markov Decision Process

8.4.1 Definition

The **Markov decision process** (MDP) is the general framework that describes the stochastic dynamic systems. It is defined as following:

Definition 8.4.1.1 (Markov Decision Process)

The **Markov decision process** is the 5-tuple $(\mathcal{S}, \mathcal{A}, p, \mathcal{R}, \gamma)$, where each elements are defined as following:

- $\mathcal{S}$ is the state space, $\mathcal{A}(s)$ is the action space associated with state s , and $\mathcal{R}(s, a)$ is a set of rewards, associated with each state-action pair (s, a) .
- $p(s', r|s, a)$ is the dynamics of the model, which is as probability of moving to state s' from s by taking action a , and obtaining reward r . Note that the state transition probability and the reward probability can be decribed with dynamics. See 8.4.1.1
- $\pi(a|s)$ is the policy, a probability of agent choosing action a in state s .
- γ is the discount rate.

with the dynamics satisfying the *Markov property*:

$$p(s_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = p(s_{t+1}|s_t, a_t) \quad (8.5)$$

$$p(r_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = p(r_{t+1}|s_t, a_t) \quad (8.6)$$

where t is the current time step.

Remark 8.4.1.1 (Transition and reward probability as dynamics)

The 4-tuple dynamics can be used to describe transition probability and reward probability as following:

$$p(s'|s, a) = \sum_{r \in \mathcal{R}(s, a)} p(s', r|s, a) \quad (8.7)$$

$$p(r|s, a) = \sum_{s' \in \mathcal{S}} p(s', r|s, a) \quad (8.8)$$

The dynamics can be either stationary or non-stationary. If the dynamics changes over time, it is non-stationary; else, it is called stationary.

Chapter 9

Bellman Equation

9.1 Introduction

TODO: Introduce the Bellman Operator and Bellman Optimality Operator. The **Bellman equation** is the heart of reinforcement learning. All reinforcement learning algorithms are designed to solve this equation.

9.2 Bellman Equation

9.2.1 State-value Function

Before diving into the Bellman equation, we introduce an important concept called the **value of a state**.

Definition 9.2.1.1 (State Value)

Let the agent follows the policy π . The **state value** $v_\pi(s)$ of state s is defined as the expected return:

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s] \quad (9.1)$$

where G_t is the discounted return $G_t = \sum_{k=t+1}^{\infty} \gamma^{k-t-1} R_k$. The function v_π itself is called **state-value function**.

State values are used to evaluate the policy π . Policies that generate greater state values are better. Then, how can we obtain te state values? The answer is to solve the Bellman equation.

9.2.2 Bellman Equation for State-Value Function

The Bellman equation is simply a set of linear equations that describe the relationships of each state values. Bellman equation can be obtained by using the recursive property of return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots \quad (9.2)$$

$$= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4}) + \dots \quad (9.3)$$

$$= R_{t+1} + \gamma G_{t+1}. \quad (9.4)$$

Substituting equation 9.1 with the recursive relationship, we obtain

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s] \quad (9.5)$$

$$= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s] \quad (9.6)$$

$$= \mathbb{E}[R_{t+1} | S_t = s] + \gamma \mathbb{E}[G_{t+1} | S_t = s] \quad (9.7)$$

Using the law of total expectation 11.2, we get

$$\mathbb{E}[R_{t+1} | S_t = s] = \sum_{a \in \mathcal{A}(s)} \pi(a|s) \mathbb{E}[R_{t+1} | S_t = s, A_t = a] \quad (9.8)$$

$$= \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{r \in \mathcal{R}(s,a)} p(r|s,a) r \quad (9.9)$$

and

$$\mathbb{E}[G_{t+1} | S_t = s] = \sum_{s' \in \mathcal{S}} \mathbb{E}[G_{t+1} | S_{t+1} = s', S_t = s] p(s'|s) \quad (9.10)$$

$$= \sum_{s' \in \mathcal{S}} \mathbb{E}[G_{t+1} | S_{t+1} = s'] p(s'|s) \quad (\text{Markov property}) \quad (9.11)$$

$$= \sum_{s' \in \mathcal{S}} v_\pi(s') p(s'|s) \quad (9.12)$$

$$= \sum_{s' \in \mathcal{S}} v_\pi(s') \sum_{a \in \mathcal{A}(s)} p(s'|s,a). \quad (9.13)$$

The final result is given in the following box:

Theorem 9.2.2.1 (Bellman Equation for State-Value Function)

Let π be a policy, and let v_π be a state-value function. Then, the following Bellman equation is satisfied:

$$v_\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}(s,a)} p(s', r | s, a) [r + \gamma v_\pi(s')] \quad (9.14)$$

9.2.3 Matrix-Vector form of Bellman Equation

We can reformulate Bellman equation in terms of matrix-vector form. For this, we rewrite the Bellman equation given in Theorem 9.2.2.1 as

$$v_\pi(s) = r_\pi(s) + \gamma \sum_{s' \in \mathcal{S}} p_\pi(s'|s) v_\pi(s'), \quad (9.15)$$

where

$$r_\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{r \in \mathcal{R}(s,a)} p(r|s,a) r \quad (9.16)$$

$$p_\pi(s'|s) = \sum_{a \in \mathcal{A}(s)} \pi(a|s) p(s'|s,a). \quad (9.17)$$

The r_π and p_π denotes the mean of immediate rewards and probability of transitioning from s to s' under policy π , respectively. Suppose that we indexed all states as s_i with $i = 1, \dots, n$, where $n = |\mathcal{S}|$. Let

$$v_\pi = \begin{bmatrix} v_\pi(s_1) \\ \vdots \\ v_\pi(s_n) \end{bmatrix}, r_\pi = \begin{bmatrix} r_\pi(s_1) \\ \vdots \\ r_\pi(s_n) \end{bmatrix}, P_\pi = \begin{bmatrix} p_\pi(s_1|s_1) & p_\pi(s_2|s_1) & \dots & p_\pi(s_n|s_1) \\ p_\pi(s_1|s_2) & p_\pi(s_2|s_2) & \dots & p_\pi(s_n|s_2) \\ \vdots & \vdots & & \vdots \\ p_\pi(s_1|s_n) & p_\pi(s_2|s_n) & \dots & p_\pi(s_n|s_n) \end{bmatrix}. \quad (9.18)$$

Then, Theorem 9.2.2.1 can be written in the following form:

Theorem 9.2.3.1 (Matrix-Vector form of Bellman Equation)

Let v_π , r_π and P_π be given as 9.18. Then, following is satisfied:

$$v_\pi = r_\pi + \gamma P_\pi v_\pi \quad (9.19)$$

Remark 9.2.3.1 (Properties of Matrix P_π)

Note that the matrix P_π given in equation 9.18 satisfies following properties:

- All the elements are nonnegative;
- The sum of the values of every row equals 1. Matrix with this property is called a *stochastic matrix*.

9.3 Solving The Bellman Equation

There are two methods in solving the Bellman equation. The closed-form solution, and the iterative solution.

9.3.1 Closed-form Solution

Since we have represented the Bellman equation as the matrix form, solving it is very easy.

Theorem 9.3.1.1 (Closed-form Solution of The Bellman Equation)

Let π be a given policy, and v_π , r_π , P_π be given as 9.18. The solution to the equation $v_\pi = r_\pi + \gamma P_\pi v_\pi$ is given as

$$v_\pi = (I - \gamma P_\pi)^{-1} r_\pi \quad (9.20)$$

Proof. Manipulate the equation as following:

$$v_\pi = r_\pi + \gamma P_\pi v_\pi \quad (9.21)$$

$$v_\pi - \gamma P_\pi v_\pi = r_\pi \quad (9.22)$$

$$(I - \gamma P_\pi) v_\pi = r_\pi. \quad (9.23)$$

Now, what we have to show is that the matrix $I - \gamma P_\pi$ is invertible. We use the fact that A is invertible $\Leftrightarrow$ eigenvalues of A are nonzero. To show this, we use Geshgorin Circle Theorem 12.2.0.1. Note that the i -th Geshgorin circles has a center at $1 - \gamma p(s_i|s_i)$, with radius $\sum_{j \neq i} \gamma p(s_j|s_i)$. Since P_π is stochastic matrix, we know that $\sum_j \gamma p(s_j|s_i) = \gamma < 1$. That is,

$\sum_{j \neq i} \gamma p_{\pi}(s_j | s_i) < 1 - \gamma p_{\pi}(s_i | s_i)$, which means that the i -th Geshgorin circle cannot contain the origin. According to the Geshgorin Circle Theorem, all the eigenvalues cannot be zero. $\square$

TODO: add properties of $I - \gamma P_{\pi}$.

9.3.2 Iterative Solution

The direct method for solving the Bellman equation is not used very well in practice, since the inverse matrix computation is very costly. Instead of the close-form method, in practice, we use iterative method to solve the Bellman equation.

Theorem 9.3.2.1 (Iterative Solution of The Bellman Equation)

Let π be a policy and v_{π} , r_{π} , and P_{π} be that of 9.18. Then, the solution of 9.2.3.1 equals the limit of following sequence $\{v_k\}$:

$$v_{k+1} = r_{\pi} + \gamma P_{\pi} v_k \quad (9.24)$$

Proof. What we have to show is that $v_k \rightarrow v_{\pi}$ as $k \rightarrow \infty$. For this, we let $\delta_k = v_k - v_{\pi}$ and show that $\delta_k \rightarrow 0$. Substituting $v_{k+1} = \delta_{k+1} + v_{\pi}$ and $v_k = \delta_k + v_{\pi}$ into equation 9.24, we have

$$\delta_{k+1} = \gamma P_{\pi} \delta_k. \quad (9.25)$$

Since P_{π} is a stochastic matrix and $0 < \gamma < 1$, we know that $\delta_k \rightarrow 0$ as $k \rightarrow \infty$. $\square$

9.4 Bellman Equation for Action-Value Function

Now we introduce the **action value**, which denotes the value of an action at a state.

9.4.1 Action Value

The action value is defined as following:

Definition 9.4.1.1 (Action Value)

Let π be a given policy. Then, the **action value** of a state-action pair is defined as

$$q_{\pi}(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]. \quad (9.26)$$

Remark 9.4.1.1 (Properties of Action Value)

Action value $q_{\pi}(s, a)$ has following properties.

- $v_{\pi} = \sum_{a \in \mathcal{A}(s)} \pi(a | s) q_{\pi}(s, a)$
- $q_{\pi}(s, a) = \sum_{r \in \mathcal{R}(s, a)} p(r | s, a) r + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) v_{\pi}(s')$

Proof. • Use the law of total expectation 11.2.

- Use the result above and [9.2.2.1](#).

□

9.4.2 Bellman Equation for Action-Value Function

The Bellman equation for action-value function is already in our hand, since actually the Remark [9.4.1.1](#) is what all we need.

Theorem 9.4.2.1 (Bellman Equation for Action-value Function)

Let π be a given policy, and let $q_\pi(s, a)$ be an action value. Then, the following Bellman equation is satisfied:

$$q_\pi(s, a) = \sum_{r \in \mathcal{R}(s, a)} p(r|s, a)r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \sum_{a' \in \mathcal{A}(s')} \pi(a'|s') q_\pi(s', a') \quad (9.27)$$

$$= \sum_{r \in \mathcal{R}(s, a)} p(r|s, a)r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_\pi(s') \quad (9.28)$$

Moreover, the matrix-vector form is given as

$$q_\pi = \tilde{r} + \gamma \tilde{P} \Pi q_\pi \quad (9.29)$$

where $[q_\pi]_{(s, a)} = q_\pi(s, a)$, $\tilde{r}_{(s, a)} = \sum_{r \in \mathcal{R}(s, a)} p(r|s, a)r$, $[\tilde{P}]_{(s, a), s'} = p(s'|s, a)$ and Π is a block diagonal matrix such that $\Pi_{s', (s', a')} = \pi(s'|a')$ for $a' \in \mathcal{A}(s')$ and other entries are all zero.

Proof. Use the Remark [9.4.1.1](#). For the matrix-vector form, try to derive it, as we did in Theorem [9.2.3.1](#).

□

Chapter 10

Bellman Optimality Equation

10.1 Introduction

This chapter introduces the Bellman optimal equation (BOE), and proves several properties of it. This section mainly follows the book [Zha25].

10.2 Optimal Policy

The reinforcement learning algorithms are all designed to find an optimal policy for given environment. Then, what is an 'optimal' policy? The definition is given below.

Definition 10.2.0.1 (Optimal Policy and Optimal State Value)

A policy π^* is optimal if and only if $v_{\pi^*}(s) \geq v_{\pi}(s)$ for all $s \in \mathcal{S}$ and for any other policy π . The state values of π^* are the optimal state values.

Definition 10.2.0.2 (Better Policy)

A policy π_1 is **better** than a policy π_2 if and only if $v_{\pi_1}(s) \geq v_{\pi_2}(s)$ for all $s \in \mathcal{S}$.

The questions that naturally arise are:

- Existence: Does the optimal policy exist?
- Uniqueness: Is the optimal policy unique?
- Stochasticity: Is the optimal policy stochastic or deterministic?
- Algorithm: How do we obtain the optimal policy and the optimal state values?

10.3 Bellman Optimality Equation

The optimal policy and optimal state values can be analysed via the **Bellman optimality equation** (BOE). The optimal state values satisfies following Bellman optimality equation. This claim in fact has two logically independent statements:

1. The BOE, viewed purely as a fixed-point equation, has a unique solution (Section 10.4);
2. The unique solution of BOE concides with the optimal stat-value function v^* , and the derived greedy policy is optimal (Section 10.5)

Note that the Section 10.4 does not assume that the state-value function is optimal.

Definition 10.3.0.1 (Bellman Optimality Equation)

Let Π be a set of possible policies and let $\gamma \in (0, 1)$ be a discount rate. Then, the Bellman Optimality Equation is defined as

$$v(s) = \max_{\pi \in \Pi} \sum_{a \in \mathcal{A}(s)} \pi(a|s) \left(\sum_{r \in \mathcal{R}} p(r|s, a)r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a)v(s') \right) \quad (10.1)$$

$$= \max_{\pi \in \Pi} \sum_{a \in \mathcal{A}(s)} \pi(a|s)q(s, a) \quad (10.2)$$

where $v(s), v(s')$ are unkown variables to be solved. Additionally, the matrix-vector form of Bellman optimality equation is given as

$$v = \max_{\pi \in \Pi} (r_{\pi} + \gamma P_{\pi} v) \quad (10.3)$$

where r_{π} and P_{π} are given as 9.18.

Using the matrix-vector form of BOE, we can express BOE very simply as following:

$$v = f(v) \quad (10.4)$$

where $f(v) = \max_{\pi \in \Pi} (r_{\pi} + \gamma P_{\pi} v)$.

10.4 Solving The Bellman Optimality Equation

10.4.1 Contraction Mapping Theorem

We propose a **contraction mapping theorem**, the main tool for solving the BOE. The contraction mapping theorem is a powerful tool for analyzing general nonlinear equations.

Definition 10.4.1.1 (Contraction Mapping)

Let $f : \mathbb{R}^d \rightarrow \mathbb{R}^d$. Then, the function f is a **contraction mapping** if there exists $\gamma \in (0, 1)$ such that

$$\|f(x_1) - f(x_2)\| \leq \gamma \|x_1 - x_2\| \quad (10.5)$$

for all $x_1, x_2 \in \mathbb{R}^d$.

Now we present the contraction mapping theorem.

Theorem 10.4.1.1 (Contraction Mapping Theorem)

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ be a contraction mapping. Then, the following properties hold for equation $x = f(x)$:

- Existence: There exists a fixed point x^* satisfying $f(x^*) = x^*$;
- Uniqueness: The fixed point x^* is unique;
- Algorithm: The following sequence $\{x_k\}$ converges to the fixed point x^* exponentially, for any initial x_0 .

$$x_{k+1} = f(x_k) \tag{10.6}$$

Proof. Part 1. (Convergence) The sequence $\{x_k\}$ defined as $x_{k+1} = f(x_k)$ converges.

We show that the sequence $\{x_k\}$ is a Cauchy sequence; i.e., there always exists an $N \in \mathcal{N}$ such that $\|x_m - x_n\| < \epsilon$ for any $\epsilon > 0$. The Cauchy sequence has a nice property; the Cauchy sequence always converges. This can be proved by using the properties of $\mathcal{R}$, that the $\mathcal{R}$ is a complete set. We won't show this here. The following holds:

$$\|x_{k+1} - x_k\| = \|f(x_k) - f(x_{k-1})\| \tag{10.7}$$

$$\leq \gamma \|x_k - x_{k-1}\| \tag{10.8}$$

$$\leq \gamma^2 \|x_{k-1} - x_{k-2}\| \tag{10.9}$$

$$\leq \gamma^3 \|x_{k-2} - x_{k-3}\| \tag{10.10}$$

$$\vdots \tag{10.11}$$

$$\leq \gamma^k \|x_1 - x_0\| \tag{10.12}$$

Using this, we can show that

$$\|x_m - x_n\| \leq \|x_m - x_{m-1}\| + \|x_{m-1} - x_{m-2}\| + \cdots + \|x_{n+1} - x_n\| \tag{10.13}$$

$$\leq \gamma^{m-1} \|x_1 - x_0\| + \gamma^{m-2} \|x_1 - x_0\| + \cdots + \gamma^n \|x_1 - x_0\| \tag{10.14}$$

$$= \|x_1 - x_0\| (\gamma^{m-1} + \gamma^{m-2} + \cdots + \gamma^n) \tag{10.15}$$

$$= \|x_1 - x_0\| \frac{\gamma^n (1 - \gamma^{m-n})}{1 - \gamma} \tag{10.16}$$

$$\leq \|x_1 - x_0\| \frac{\gamma^n}{1 - \gamma} \tag{10.17}$$

where $m \geq n$, without loss of generality. Since $\gamma \in (0, 1)$, the last line converges to 0. That is, the sequence $\{x_k\}$ is a Cauchy sequence.

Part 2. (Existence) $\lim_{k \rightarrow \infty} x_k$ is a solution of fixed-point equation $x = f(x)$.

Let $x^* = \lim_{k \rightarrow \infty} x_k$. Then, following holds.

$$\|x^* - f(x^*)\| \leq \|x^* - x_k\| + \|f(x_k) - f(x^*)\| + \|x_k - f(x_k)\| \tag{10.18}$$

$$\leq \|x^* - x_k\| + \gamma \|x_k - x^*\| + \|x_k - f(x_k)\| \tag{10.19}$$

Using the results of *Part 1* and squeeze theorem, we can show that $x^* = f(x^*)$.

Part 3. (Uniqueness) The fixed-point x^* is unique.

Suppose there is another fixed point x' . Then, $\|x^* - x'\| = \|f(x^*) - f(x')\| \leq \gamma\|x^* - x'\|$. Since $\gamma \in (0, 1)$, the only possibility is $x^* = x'$.

□

10.4.2 Contraction Property of Bellman Optimality Equation

Now, to use the contraction mapping theorem, we prove that the right-hand side of BOE 10.3 is a contraction mapping.

Theorem 10.4.2.1 (Contraction Property of Bellman Optimality Equation)

Let the function f be that of equation 10.3. Then, the following contraction mapping property holds for any $v_1, v_2 \in \mathbb{R}^{|S|}$:

$$\|f(v_1) - f(v_2)\|_\infty \leq \gamma\|v_1 - v_2\|_\infty \quad (10.20)$$

where $\gamma \in (0, 1)$ is the discount rate, and $\|\cdot\|_\infty$ is the maximum norm, which is the maximum absolute value of the elements in a vector.

Proof. Let $v_1, v_2 \in \mathbb{R}^S$, and let $\pi_1^* = \arg \max_\pi (r_\pi + \gamma P_\pi v_1)$, $\pi_2^* = \arg \max_\pi (r_\pi + \gamma P_\pi v_2)$. Then,

$$f(v_1) = \max_\pi (r_\pi + \gamma P_\pi v_1) = r_{\pi_1^*} + \gamma P_{\pi_1^*} v_1 \geq r_{\pi_2^*} + \gamma P_{\pi_2^*} v_1 \quad (10.21)$$

$$f(v_2) = \max_\pi (r_\pi + \gamma P_\pi v_2) = r_{\pi_2^*} + \gamma P_{\pi_2^*} v_2 \geq r_{\pi_1^*} + \gamma P_{\pi_1^*} v_2 \quad (10.22)$$

where $\geq$ is applied elementwise. This result leads to

$$f(v_1) - f(v_2) \leq \gamma P_{\pi_1^*} (v_1 - v_2) \quad (10.23)$$

$$f(v_2) - f(v_1) \leq \gamma P_{\pi_2^*} (v_2 - v_1) \quad (10.24)$$

Let $z = \max \{|\gamma P_{\pi_1^*} (v_1 - v_2)|, |\gamma P_{\pi_2^*} (v_2 - v_1)|\}$, where $\max(\cdot)$ and $|\cdot|$ are all applied elementwise. Then,

$$\|f(v_1) - f(v_2)\|_\infty \leq \|z\|_\infty. \quad (10.25)$$

Now, let z_i be the i -th entry of z , and let $p_i^\top$ and $q_i^\top$ be the i -th row of $P_{\pi_1^*}$ and $P_{\pi_2^*}$, respectively. Then, $z_i = \max \{\gamma |p_i^\top (v_1 - v_2)|, \gamma |q_i^\top (v_2 - v_1)|\}$. Since all the elements of vector p_i is nonnegative and thier sum equals 1,

$$|p_i^\top (v_1 - v_2)| \leq p_i^\top |v_1 - v_2| \leq \|v_1 - v_2\|_\infty. \quad (10.26)$$

where $|\cdot|$ is applied elementwise. Similarly, $|q_i^\top (v_2 - v_1)| \leq \|v_2 - v_1\|_\infty$. Therefore $\|z\|_\infty \leq \gamma\|v_1 - v_2\|_\infty$. Substituting to 10.25,

$$\|f(v_1) - f(v_2)\|_\infty \leq \gamma\|v_1 - v_2\|_\infty. \quad (10.27)$$

□

10.5 Optimal Policy, Optimal State Value and Bellman Optimality Equation

We have a following Theorem of BOE:

Theorem 10.5.0.1 (Solution of the Bellman Optimality Equation)

Let f be as the equation 10.3. Then, the BOE $v = f(v)$ has following properties:

1. There exists a unique solution;
2. The solution can be obtained iteratively by $v_{k+1} = f(v_k)$.

Proof. All the previous Theorems (10.4.1.1, 10.4.2.1) proves this trivially. □

Now, we should prove that the solution of the Bellman optimality equation is really an optimal state-value function, and prove that the greedy policy derived from it is the optimal policy.

Theorem 10.5.0.2 (Optimality of v^* and π^*)

Let v^* be the solution of BOE $v = f(v)$, and let $\pi^* = \arg \max_{\pi \in \Pi} (r_\pi + \gamma P_\pi v^*)$. The, the v^* is an optimal state value function, and the π^* is an optimal policy. That is, for any policy π , following holds;

$$v^* = v_{\pi^*} \geq v_\pi \quad (10.28)$$

where v_π is the state value function of π , and $\geq$ is applied elementwise.

Proof. For any policy π , $v_\pi = r_\pi + \gamma P_\pi v_\pi$ holds (Bellman equation). Then,

$$v^* = \max_{\tau \in \Pi} (r_\tau + \gamma P_\tau v^*) \quad (10.29)$$

$$= r_{\pi^*} + \gamma P_{\pi^*} v^* \quad (10.30)$$

$$\geq r_\pi + \gamma P_\pi v^* \quad (10.31)$$

and we have

$$v^* - v_\pi \geq \gamma P_\pi (v^* - v_\pi). \quad (10.32)$$

Using this inequality repeatedly, we have $v^* - v_\pi \geq \gamma^n P_\pi^n (v^* - v_\pi)$ for any $n \in \mathbb{N}$. Since $\gamma \in (0, 1)$ and P_π is a nonnegative matrix with all its elements less than or equal to 1, $\lim_{n \rightarrow \infty} \gamma^n P_\pi^n (v^* - v_\pi) = 0$. That is, $v^* - v_\pi \geq \lim_{n \rightarrow \infty} \gamma^n P_\pi^n (v^* - v_\pi) = 0$. □

The following theorem examines the optimal policy more precisely.

Theorem 10.5.0.3 (Greedy Optimal Policy)

For any $s \in \mathcal{S}$, the following deterministic policy is an optimal policy for solving the BOE.

$$\pi^*(a|s) = \begin{cases} 1, & a = \arg \max_a q^*(a, s) \\ 0, & a \neq \arg \max_a q^*(a, s) \end{cases} \quad (10.33)$$

where $q^*(s, a) = \sum_{r \in \mathcal{R}} p(r|s, a)r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a)v^*(s')$.

Proof. From Theorem 10.5.0.1 and 10.3.0.1, we see that

$$v^*(s) = \max_{\pi \in \Pi} \sum_{a \in \mathcal{A}} \pi(a|s) q^*(s, a) \quad (10.34)$$

$$\leq \left(\max_{\pi \in \Pi} \sum_{a \in \mathcal{A}} \pi(a|s) \right) \max_{a'} q^*(s, a') \quad (10.35)$$

$$= \max_{a'} q^*(s, a') \quad (10.36)$$

and the given deterministic policy π^* satisfies the equality of inequality 10.35. $\square$

The important truth is that the optimal state value is unique but the optimal policy is not. The Theorem 10.5.0.3 only shows that there always exists at least one deterministic optimal policy. There still can be other optimal policies, too, stochastic or whatever.

10.6 Factors that Influence Optimal Policies

Now, we analyze the factors of BOE that influences optimal policies. The factors are: reward and discount rate.

10.6.1 Reward

10.6.2 Discount Rate

Part III

Appendices

Chapter 11

Probability Theory

11.1 Introduction

A short, useful theorems of probability theory.

11.2 Law of Total Expectation

The followinga are the law of total expectations.

$$\mathbb{E}[X] = \sum_a \mathbb{E}[X|A = a]p(a) \quad (11.1)$$

$$\mathbb{E}[X|A = a] = \sum_b \mathbb{E}[X|A = a, B = b]p(b|a) \quad (11.2)$$

Chapter 12

Linear Algebra

12.1 Introduction

A short useful theorems of linear algebra.

12.2 Geshgorin Circle Theorem

Theorem 12.2.0.1 (Geshgorin Circle Theorem)

Bibliography

- [Gra20] Laura Graesser. *Foundations of deep reinforcement learning. Theory and practice in Python*. Ed. by Wah Loon Keng. Addison Wesley data & analytics series. Includes bibliographical references and index. Boston: Addison-Wesley, 2020. 379 pp. ISBN: 0135172381.
- [Sut20] Richard S. Sutton. *Reinforcement learning. An introduction*. Ed. by Andrew Barto. Second edition. Adaptive computation and machine learning. Hier auch später erschienene, unveränderte Nachdrucke. Cambridge, Massachusetts: The MIT Press, 2020. 526 pp. ISBN: 9780262039246.
- [Zha25] Shiyu Zhao. *Mathematical foundations of reinforcement learning*. [Tsinghua]: Tsinghua University Press, 2025. 1275 pp. ISBN: 9789819739448.